

Discrete Planning


Vikas Dhiman

Monday 8th September, 2025

1 Required readings

- Ch 7 and 23, [Carmen's Intro to Algorithms](#)
- [Breadth-first search and its uses \(article\) | Khan Academy](#)
- [PythonRobotics/dijkstra.py](#)
- [Discrete Planning by Laval](#)
- 3.5.2 of [Russel and Norvig: Artificial Intelligence](#)

2 Discrete planning in Robotics



2.0.1 Abstraction of a planning problem

1. State space $s \in \mathcal{S}$. For example, 2D coordinate of a grid $s = (x, y)$.
2. Action space per state $u \in \mathcal{U}(s)$. For example, up, down, left right movement can be encoded as $\mathcal{U}(s_t) = \{(0, -1), (0, 1), (1, 0), (-1, 0)\}$.

State = ? / State space
 ↳ a) Time varying property

b) Future action should be completely described by state

↳ $a_t \perp s_{t-1} | s_t$.
Markovian


For a robot that cannot move sideways and have to turn
 state = (x, y, θ)
 $(-\pi, \pi)$

x, y, z / θ, ϕ, ψ
 3D / roll/pitch/yaw
 ↑ Euler angles

$IP(s_{t+1} | s_t, u_t) = ?$
 state transition probabilities

system dynamics

$s_{t+1} = f(s_t, u_t)$
 ↑ next state ↑ state action



$s_t = \begin{bmatrix} x_t \\ y_t \end{bmatrix}$
 $u_t = \begin{bmatrix} \Delta x_t \\ \Delta y_t \end{bmatrix}$
 down ↑ up
 left ↑ right

3. State transition function $s_{t+1} = f(s_t, u_t)$. For example, the up-down-left-right action can be combined as addition to get the next state $s_{t+1} = s_t + u_t$.
4. Initial State $s_I \in S$
5. Goal states $s_G \subseteq S$

- ① State Space : all possible states
- ② Action space : all possible action
- ③ $s_{t+1} = f(s_t, u_t)$
- ④ s_0
- ⑤ Goal state s_g


2.0.2 A Graph

A graph $G = \{V, E\}$ is defined by a set of vertices V and a set of edges E such that each edge $e \in E$ is formed by a pair of start and end vertices $e = (v_s, v_e), v_s \in V, v_e \in V$. The first vertex is called the start of the edge $v_s = \text{start}(e)$ and second vertex is called the end $v_e = \text{end}(e)$.

A discrete planning problem can be converted into a graph by defining

1. Vertices as the state space $V = S$.
2. The action space at each state as the edges connected to that vertex/state, $U(s_t) = \{(s_t, s_j) \mid (s_t, s_j) \in E\}$ set of edge
3. State transition function is the other end of the edge, $s_{t+1} = f(s_t, u_t) = \text{end}(u_t)$, where $s_t = \text{start}(u_t)$.

$s_{t+1} = s_t + u_t$
 $f(s_t, u_t)$
 $s_{t+1} = \begin{cases} [x+1, y] & \text{if } u_t = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \\ [x-1, y] & \text{if } u_t = \begin{bmatrix} -1 \\ 0 \end{bmatrix} \\ [x, y+1] & \text{if } u_t = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \\ [x, y-1] & \text{if } u_t = \begin{bmatrix} 0 \\ -1 \end{bmatrix} \end{cases}$



2.0.3 Representations of Graphs

(Chapter 23 of Introduction to Algorithms by Carmen et al)

- ① Adjacency list
- ② Adjacency matrix
- ③ Edge list





Figure 23.1 (c) The adjacency-matrix representation of G.

	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0


Undirected graph

Programmatically you can represent a adjacency list as python lists
 # Python lists are not linked lists, they are arrays under the hood.

```
G_adjacency_list = {
    key: 1 : [2, 5], # node : list of neighbors as values
        2 : [1, 5, 3, 4],
        3 : [2, 4],
        4 : [2, 5, 3],
        5 : [4, 1, 2]
}
```

Prefer to represent a matrix in python either as a list of lists or a numpy array

Discrete planning



$\{$
 $1: [(2, 2.0), (5, 3.0)],$
 $2: [(1, 2.0), (3, 3.5), (4, 2.5), (5, 1.5)],$
 $\vdots$
 $\}$
 $O(|E|/|V|)$


$\begin{matrix} & 1 & 2 & 3 & 4 & 5 \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{pmatrix} 0 & 2.0 & 0 & 0 & 3.0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{pmatrix} \end{matrix}$

$(2, 4)?$
 $O(1)$

```


import numpy as np
G_adjacency_matrix = np.array([
    [0, 1, 0, 0, 1],
    [1, 0, 1, 1, 1],
    [0, 1, 0, 1, 0],
    [0, 1, 1, 0, 1],
    [1, 1, 0, 1, 0]
])

```



$\mathcal{E} =$ # Edge list is another possible representation
 $G_edge_list = \{$
 $(1, 2), (1, 5),$
 $(2, 1), (2, 5), (2, 3), (2, 4),$
 $(3, 2), (3, 4),$
 $(4, 2), (4, 5), (4, 3),$
 $(5, 4), (5, 1), (5, 2)$
 $\}$

$\{$
 $(1, 2): 3.0$
 $(1, 5): 4.5$
 $\}$



Directed graph representation

```

# Programmatically you can represent a adjacency list as python lists
# Python lists are not linked lists, they are arrays under the hood.
G_adjacency_list = {
    1 : [2, 4],
    2 : [5],
    3 : [6, 5],
    4 : [2],
    5 : [4],
    6 : [6]
}

```

```

# Prefer to represent a matrix in python either as a list of lists or a numpy array
import numpy as np
G_adjacency_matrix = np.array([
    [0, 1, 0, 1, 0, 0],
    [0, 0, 0, 0, 1, 0],
    [0, 0, 0, 0, 1, 1],
    [0, 1, 0, 0, 0, 0],
    [0, 0, 0, 1, 0, 0],
    [0, 0, 0, 0, 0, 0]
])

```

```

        [0, 0, 0, 0, 0, 1]
    ])

# Edge list is another possible representation
G_edge_list = [
    (1, 2), (1, 4),
    (2, 5),
    (3, 6), (3, 5),
    (4, 2),
    (5, 6)
]

# Exercise 1

# Write a function that converts a graph in adjacency list format to adjacency matrix and vice versa
def adjacency_list_to_matrix(G_adj_list):
    G_adj_mat = None # TODO: Write code to convert to adj_mat
    return G_adj_mat

def adjacency_matrix_to_list(G_adj_mat):
    G_adj_list = None # TODO: Write code to convert to adj_list
    return G_adj_list





# Use the above graphs to test
print(adjacency_list_to_matrix(G_adjacency_list))
print(adjacency_matrix_to_list(G_adjacency_matrix))

None
None


```

Tree
vs
2.0.4 Graph Search algorithms
Traversal

Trees don't have cycles/Loops



1. Breadth First Search ✓
2. Depth First Search ✓






Figure 23.3 The operation of BFS on an undirected graph. Tree edges are shown shaded as they are produced by BFS. Within each vertex u is shown $d[u]$. The queue Q is shown at the beginning of each iteration of the while loop of lines 9–18. Vertex distances are shown next to vertices in the queue.

Breadth first search (BFS)

```
from queue import Queue, LifoQueue, PriorityQueue
```

```
graph = {
    's' : ['w', 'r'],
    'r' : ['v'],
    'w' : ['t', 'x'],
    'x' : ['y'],
    't' : ['u'],
    'u' : ['y']
}

def bfs(graph, start, debug=False):
    seen = set() # Set for seen nodes (contains both frontier and dead states)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = Queue() # Frontier of unvisited nodes as FIFO
    node2dist = {start : 0} # Keep track of distances
    search_order = []
    seen.add(start)
    frontier.put(start)

    i = 0 # step number
    while not frontier.empty():
        # Creating loop to visit each node
        if debug: print("%d) Q = " % i, list(frontier.queue), end='; ')
        if debug: print("dists = " , [node2dist[n] for n in frontier.queue])
        m = frontier.get() # Get the oldest addition to frontier
        search_order.append(m)
```

```


    for neighbor in graph.get(m, []):
        if neighbor not in seen:
            seen.add(neighbor)
            frontier.put(neighbor)
            node2dist[neighbor] = node2dist[m] + 1
        else:
            assert node2dist[neighbor] <= node2dist[m] + 1, 'this should not happen in BFS'
            node2dist[neighbor] = min(node2dist[neighbor], node2dist[m] + 1)

    i += 1
    if debug: print("%d) Q = " % i, list(frontier.queue))
    return search_order, node2dist

print("Following is the Breadth -First Search order")
print(bfs(graph, 's', debug=True))    # function calling

Following is the Breadth -First Search order
0) Q = ['s']; dists = [0]
1) Q = ['w', 'r']; dists = [1, 1]
2) Q = ['r', 't', 'x']; dists = [1, 2, 2]
3) Q = ['t', 'x', 'v']; dists = [2, 2, 2]
4) Q = ['x', 'v', 'u']; dists = [2, 2, 3]
5) Q = ['v', 'u', 'y']; dists = [2, 3, 3]
6) Q = ['u', 'y']; dists = [3, 3]
7) Q = ['y']; dists = [3]
8) Q = []
(['s', 'w', 'r', 't', 'x', 'v', 'u', 'y'], {'s': 0, 'w': 1, 'r': 1, 't': 2, 'x': 2, 'v': 2,

```



BFS
Depth first search
gray = frontier
seen = black




Figure 23.3 The operation of BFS on an undirected graph. Tree edges are shown shaded as they are produced by BFS. Within each vertex u is shown $d[u]$. The queue Q is shown at the beginning of each iteration of the while loop of lines 9–18. Vertex distances are shown next to vertices in the queue.


```

graph = {
  's' : ['w', 'r'],
  'r' : ['v'],
  'w' : ['t', 'x'],
  'x' : ['y'],
  't' : ['u'],
  'u' : ['y']
}

```


DFS
Depth first
Search


BFS



seen = set({})
frontier = Queue()
DFS: FILO stack
BFS: FIFO
nodes that have been visited but their neighbors have not been
FILO: First In Last Out
Stack




FIFO: First In First Out




find neighbors of w

seen
[s, w]
frontier (FILO)
[r, w]





seen
seen
[s, w, t, x, v]
frontier
[r, t, x]
↑
intermediate

find neighbors of x



seen
[s, w, x, t]
frontier
[r, t, y]

BFS & DFS for shortest path



Breadth first search
gives shortest path

for graphs with uniform
edge weights / distance

BFS → Dijkstra algorithm
(updates the visiting order based on
shortest distance)

FIFO
queue

Priority Queue


Out ←

4	3	3.5	2.1	2.0	1.0
v_1	v_2	v_3	v_4	v_5	v_6

← Priority

Highest priority node gets visited first


Dijkstra



seen

frontier

0
s




seen

s	u	x
---	---	---

frontier

-5	-10
x	u

Priority Queue



seen


s	x	u	v	y
---	---	---	---	---

frontier

-14	-7	-8
v	y	u

Priority Queue

-7	-8	-14
y	u	v



seen

s	x	u	v	y
---	---	---	---	---

frontier


Priority Queue

-8	-11	
u	v	

s	x	u	v	y
---	---	---	---	---

frontier

-9
v



computational complexity

LIFO Queue / stack (n)


FIFO Queue (n)

Priority Queue

enqueueing add $O(1)$ dequeueing remove $O(1)$



array



std::vector $O(1)$ $O(1)$


4	2	3	1	5
a	b	c	d	e

$O(1)$ $O(n)$

$O(n)$ $O(1)$

$O(\log(n))$ $O(1)$

heap data structure



Priority Queue → Priority Queue Updateable

}

```
def dfs(graph, start, debug=False):
    seen = set([start]) # List for seen nodes (contains both frontier and dead states)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = LifoQueue() # Frontier of unvisited nodes as FIFO
    node2dist = {start : 0} # Keep track of distances
    search_order = [] # Keep track of search order
    frontier.put(start)
```

frontier = Queue() →

```
    i = 0 # step number
    while not frontier.empty(): # Creating loop to visit each node
        if debug: print("%d Q = " % i, list(frontier.queue), end='; ')
        if debug: print("dists = " , [node2dist[n] for n in frontier.queue])
        m = frontier.get() # Get the oldest addition to frontier
        search_order.append(m)
```

```
        # adjacency list
        for neighbor in graph.get(m, []):
            if neighbor not in seen:
                seen.add(neighbor)
                frontier.put(neighbor)
                node2dist[neighbor] = node2dist[m] + 1
            else:
                node2dist[neighbor] = min(node2dist[neighbor], node2dist[m] + 1)
        i += 1
    if debug: print("%d Q = " % i, list(frontier.queue))
    return search_order, node2dist
```

graph[m] may throw Key Error
edge-distance


Driver Code

```
print("Following is the Depth -First Search path")
print(dfs(graph, 's', debug=True)) # function calling
```



Following is the Depth -First Search path

```
0) Q = ['s']; dists = [0]
1) Q = ['w', 'r']; dists = [1, 1]
2) Q = ['w', 'v']; dists = [1, 2]
3) Q = ['w']; dists = [1]
4) Q = ['t', 'x']; dists = [2, 2]
5) Q = ['t', 'y']; dists = [2, 3]
6) Q = ['t']; dists = [2]
7) Q = ['u']; dists = [3]
8) Q = []
(['s', 'r', 'v', 'w', 'x', 'y', 't', 'u'], {'s': 0, 'w': 1, 'r': 1, 'v': 2, 't': 2, 'x': 2,
```


2.0.5 Search order in BFS vs DFS vs Dijkstra




Breadth first search vs Depth first search



Breadth first search vs Dijkstra





2.0.6 Converting a maze search to a graph search

```
import matplotlib.pyplot as plt
import numpy as np
maze_str = \
"""
+++++++
  +      +
+ + + +++
+ + +   +
+ + +   +
+ + +++ +
+      + +
+ +++ + +
+      +
+++++++
"""

class Maze:
    def __init__(self, maze_str, freepath=' '):
        self.maze_lines = [l for l in maze_str.split("\n")
                           if len(l)]
        self.FREEPATH = freepath
        self.fig = None

    def get(self, node, default):
        (r, c) = node
        m_row = self.maze_lines[r]
        nbrs = []
        if c - 1 >= 0 and m_row[c - 1] == self.FREEPATH:
            nbrs.append((r, c - 1))
        if c + 1 < len(m_row) and m_row[c + 1] == self.FREEPATH:
            nbrs.append((r, c + 1))
        if r - 1 >= 0 and self.maze_lines[r - 1][c] == self.FREEPATH:
```

```


        nbrs.append((r -1, c))
        if r+1 < len(self.maze_lines) and self.maze_lines[r+1][c] == self.FREEPATH:
            nbrs.append((r+1, c))
        return nbrs if len(nbrs) else default
init_plots = init_plots
plot_maze = plot_maze
plot_step = plot_step
plot_path = plot_path
animate_search_path = animate_search_path

import matplotlib.pyplot as plt
import matplotlib.animation as animation
import matplotlib as mpl
%matplotlib inline
mpl.rc('animation', html='jshtml')

maze = Maze(maze_str)
search_path, node2dist = bfs(maze, (1, 0)) # prints the order of search all the searched nodes
maze.plot_maze()

maze.animate_search_path(search_path, node2dist)

```




```

def bfs_path(graph, start, goal):
    """
    Returns success and node2parent

    success: True if goal is found otherwise False
    node2parent: A dictionary that contains the nearest parent for node
    """
    seen = [start] # List for seen nodes.
    # Frontier is the boundary between seen and unseen
    frontier = Queue() # Frontier of unvisited nodes as FIFO
    node2parent = dict() # Keep track of nearest parent for each node (requires node to be h
    frontier.put(start)

    while not frontier.empty():          # Creating loop to visit each node
        m = frontier.get() # Get the oldest addition to frontier
        if m == goal:
            return True, node2parent

        for neighbor in graph.get(m, []):
            if neighbor not in seen:
                seen.append(neighbor)
                frontier.put(neighbor)
                node2parent[neighbor] = m
    return False, []

def backtrack_path(node2parent, start, goal):
    c = goal
    r_path = [c]
    parent = node2parent.get(c, None)
    while parent != start:
        r_path.append(parent)
        c = parent
        parent = node2parent.get(c, None) # Keep getting the parent until you reach the start
        #print(parent)
    r_path.append(start)
    return reversed(r_path) # Reverses the path

maze = Maze(maze_str)
start = (1, 0)
goal = (8, 8)
success, node2parent = bfs_path(maze, (1, 0), (8, 8))
path = backtrack_path(node2parent, (1, 0), (8, 8))
#print(list(path))
maze.plot_path(path) # Draws all the searched nodes
plt.show()
#node2parent

```





Figure 25.5 The execution of Dijkstra's algorithm. The source is the leftmost vertex. The shortest-path estimates are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. Black vertices are in the set S , and white vertices are in the priority queue $Q = V - S$. (a) The situation just before the first iteration of the **while** loop of lines 4–8. The shaded vertex has the minimum d value and is chosen as vertex u in line 5. (b)–(f) The situation after each successive iteration of the **while** loop. The shaded vertex in each part is chosen as vertex u in line 5 of the next iteration. The d and π values shown in part (f) are the final values.

Dijkstra algorithm

2.0.7 PriorityQueue

PriorityQueue returns the smallest (or the largest) item in the queue faster than other data structures

```
import itertools
```


```

class MazeD(Maze):
    def get(self, node, default):
        nbrs = Maze.get(self, node, default)
        return zip(nbrs, itertools.repeat(1))

maze = MazeD(maze_str)
success, search_path, node2parent, node2dist = dijkstra(maze, (1, 0), (8, 8))
print(success, node2parent)
if success:
    path = backtrace_path(node2parent, (1, 0), (8, 8))
    maze.plot_path(path) # Draws all the searched nodes

True {(1, 0): None, (1, 1): (1, 0), (2, 1): (1, 1), (3, 1): (2, 1), (4, 1): (3, 1), (5, 1):

```




```


maze = Maze8(maze_str)
success, search_path, node2parent, node2dist = dijkstra(maze, start_pos, goal_pos)
#print(success, search_path)
assert success
anim = maze.animate(search_path)
path = backtrace_path(node2parent, start_pos, goal_pos)
#maze.init_plots(reinit=True)
path_plot = maze.plot_path(path, color='k') # Draws the traced shortest path
anim

```

```
/tmp/ipykernel_3055857/955263672.py:37: UserWarning: frames=<generator object batched at 0x...
anim = animation.FuncAnimation(self.fig, self._plot_path,
```



```
path = backtrace_path(node2parent, (35, 9), (5, 50))
maze.init_plots(reinit=True)
maze.plot_path(path, color='r') # Draws the traced shortest path
plt.show()
```




2.1 Computational complexity notation

Asymptotically tight bound $\Theta(g(n)) = \{f(n) : \text{there exist positive constants } c_1, c_2, \text{ and } n_0 \text{ such that } 0 \leq c_1g(n) \leq f(n) \leq c_2g(n) \text{ for all } n \geq n_0\}$

Asymptotic upper bound $O(g(n)) = \{f(n) : \text{there exist positive constants } c_2, \text{ and } n_0 \text{ such that } 0 \leq f(n) \leq c_2g(n) \text{ for all } n \geq n_0\}$

Asymptotic lower bound $\Sigma(g(n)) = \{f(n) : \text{there exist positive constants } c_1, \text{ and } n_o \text{ such that } 0 \leq c_1 g(n) \leq f(n) \text{ for all } n \geq n_o\}$



2.1.1 Computational complexity of BFS

```
# Write down the computational complexity of each line in big -O notation O()
# Assume the graph has |V| nodes and |E| edges
def bfs_barebones(graph, start):
    seen = {start} # Set for seen nodes (contains both frontier and dead states) # O(1)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = Queue() # Frontier of unvisited nodes as FIFO # O(1)
    frontier.put(start) # O(1)

    while not frontier.empty(): # Creating loop to visit each node # O(|V|)
        m = frontier.get() # Get the oldest addition to frontier # O(|V| * 1)

        for neighbor in graph.get(m, []): # O(|V| * |E|/|V|) = O(|E|)
            if neighbor not in seen: # O(|E| * 1)
                seen.add(neighbor) # O(|E| * 1)
                frontier.put(neighbor) # O(|E| * 1)

# The computational complexity of BFS is O(|E|). Some books write it as O(|V| + |E|)
# where O(|V|) is the cost of initializing states of different nodes
```

2.1.2 Computational complexity of Dijkstra

```
# Write down the computational complexity of each line in big -O notation O()
# Assume the graph has |V| nodes and |E| edges
def dijkstra_barebones(graph, start):
    seen = {start} # Set for seen nodes (contains both frontier and dead states) # O(1)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = PriorityQueue() # Frontier of unvisited nodes as PriorityQueue # O(1)
    frontier.put(PItem(0, start)) # O(1)
    node2dist = {start: 0} # Keep track of cost to arrive at each node # O(1)

    while not frontier.empty(): # Creating loop to visit each node
        dist_and_node = frontier.get() # Get the smallest dist node # O(|V| * log(|V|))
        m_dist = dist_and_node.dist
        m = dist_and_node.node
```

```

    for neighbor, edge_dist in graph.get(m, []): #  $O(|V| * |E|/|V|) = O(|E|)$ 
        if neighbor not in seen: #  $O(|E| * 1)$ 
            seen.add(neighbor) #  $O(|E| * 1)$ 
            frontier.put(neighbor) # #  $O(|E| * \log(1))$  # for fibonacci heap
            node2dist[neighbor] = m_dist + edge_dist #  $O(1)$ 
        elif node2dist[neighbor] > m_dist + edge_dist: #  $O(1)$ 
            node2dist[neighbor] = m_dist + edge_dist #  $O(1)$ 

# The computational complexity of Dijkstra is  $O(|V|\log(|V|) + |E|)$  when implemented
# using a Fibonacci heap based PriorityQueue

```

2.2 PriorityQueue (Heaps Chapter 7 of Carmen's intro to algorithms)




Figure 7.1 A heap viewed as (a) a binary tree and (b) an array. The number within the circle at each node in the tree is the value stored at that node. The number next to a node is the corresponding index in the array.

Heap property

1. $H[\text{Parent}(i)] \geq H[i]$
2. $\text{Parent}(i) = \text{ceil}(i/2)$
3. $\text{LeftChild}(i) = 2i$
4. $\text{RightChild}(i) = 2i+1$




Figure 7.2 The action of $\text{HEAPIFY}(A, 2)$, where $\text{heap-size}[A] = 10$. (a) The initial configuration of the heap, with $A[2]$ at node $i = 2$ violating the heap property since it is not larger than both children. The heap property is restored for node 2 in (b) by exchanging $A[2]$ with $A[4]$, which destroys the heap property for node 4. The recursive call $\text{HEAPIFY}(A, 4)$ now sets $i = 4$. After swapping $A[4]$ with $A[9]$, as shown in (c), node 4 is fixed up, and the recursive call $\text{HEAPIFY}(A, 9)$ yields no further change to the data structure.

Heapify

```

HEAPIFY( $A, i$ )
1   $l \leftarrow \text{LEFT}(i)$ 
2   $r \leftarrow \text{RIGHT}(i)$ 
3  if  $l \leq \text{heap-size}[A]$  and  $A[l] > A[i]$ 
4      then  $\text{largest} \leftarrow l$ 
5      else  $\text{largest} \leftarrow i$ 
6  if  $r \leq \text{heap-size}[A]$  and  $A[r] > A[\text{largest}]$ 
7      then  $\text{largest} \leftarrow r$ 
8  if  $\text{largest} \neq i$ 
9      then exchange  $A[i] \leftrightarrow A[\text{largest}]$ 
10     HEAPIFY( $A, \text{largest}$ )

```

Heapify pseudocode




Figure 7.5 The operation of HEAP-INSERT. (a) The heap of Figure 7.4(a) before we insert a node with key 15. (b) A new leaf is added to the tree. (c) Values on the path from the new leaf to the root are copied down until a place for the key 15 is found. (d) The key 15 is inserted.


Heap Insert

Operation	find-min	delete-min	insert	decrease-key	meld
Binary ^[9]	$\Theta(1)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$\Theta(n)$
Leftist	$\Theta(1)$	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)$	$\Theta(\log n)$
Binomial ^{[9][10]}	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)^{[a]}$	$\Theta(\log n)$	$O(\log n)$
Skew binomial ^[11]	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(\log n)$	$O(\log n)^{[b]}$
Pairing ^[12]	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$o(\log n)^{[a][c]}$	$\Theta(1)$
Rank-pairing ^[15]	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$\Theta(1)^{[a]}$	$\Theta(1)$
Fibonacci ^{[9][2]}	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$\Theta(1)^{[a]}$	$\Theta(1)$
Strict Fibonacci ^[16]	$\Theta(1)$	$O(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
Brodal ^{[17][d]}	$\Theta(1)$	$O(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
2-3 heap ^[19]	$O(\log n)$	$O(\log n)^{[a]}$	$O(\log n)^{[a]}$	$\Theta(1)$	?

Heap runtimes

2.3 A-star (A*) algorithm

(Required reading: 3.5.2 of Russel and Norvig: Artificial Intelligence)



```
from hw2_solution import PriorityQueueUpdatable
```



```

import sys
def astar(graph, heuristic_dist_fn, start, goal, debug=False, debugf=sys.stdout):
    """
    edgecost: cost of traversing each edge

    Returns success and node2parent

    success: True if goal is found otherwise False
    node2parent: A dictionary that contains the nearest parent for node
    """
    seen = set([start]) # Set for seen nodes.
    # Frontier is the boundary between seen and unseen
    frontier = PriorityQueueUpdatable() # Frontier of unvisited nodes as a Priority Queue
    node2parent = {start : None} # Keep track of nearest parent for each node (requires node2parent)
    hfn = heuristic_dist_fn # make the name shorter
    node2dist = {start: 0 } # Keep track of cost to arrive at each node
    search_order = []
    frontier.put(PItem(0 + hfn(start, goal), start)) # < - - - - - Different from dijkstra

    if debug: debugf.write("goal = " + str(goal) + '\n')
    i = 0
    while not frontier.empty():
        # Creating loop to visit each node
        dist_m = frontier.get() # Get the smallest addition to the frontier
        if debug: debugf.write("%d) Q = " % i + str(list(frontier.queue)) + '\n')
        if debug: debugf.write("%d) node = " % i + str(dist_m) + '\n')
        #if debug: print("dists = " , [node2dist[n.node] for n in frontier.queue])
        m = dist_m.node
        m_dist = node2dist[m]
        search_order.append(m)
        if goal is not None and m == goal:
            return True, search_order, node2parent, node2dist

        for neighbor, edge_cost in graph.get(m, []):
            old_dist = node2dist.get(neighbor, float("inf"))
            new_dist = edge_cost + m_dist
            if neighbor not in seen:
                seen.add(neighbor)
                frontier.put(PItem(new_dist + hfn(neighbor, goal), neighbor)) # < - - - - - Different from dijkstra
                node2parent[neighbor] = m
                node2dist[neighbor] = new_dist
            elif new_dist < old_dist:
                node2parent[neighbor] = m
                node2dist[neighbor] = new_dist
            # ideally you would update the dist of this item in the priority queue
            # as well. But python priority queue does not support fast updates
            # - - - - - Different from dijkstra - - - - -

```

```


        old_item = PItem(old_dist + hfn(neighbor, goal), neighbor)
        if old_item in frontier:
            frontier.replace(
                old_item,
                PItem(new_dist + hfn(neighbor, goal), neighbor))
        i += 1
    if goal is not None:
        return False, [], {}, node2dist
    else:
        return True, search_order, node2parent, node2dist

import math
from functools import partial

def euclidean_heurist_dist(node, goal, scale=1):
    x_n, y_n = node
    x_g, y_g = goal
    return scale*math.sqrt((x_n - x_g)**2 + (y_n - y_g)**2)

/tmp/ipykernel_3055857/955263672.py:37: UserWarning: frames=<generator object batched at 0x7...
    anim = animation.FuncAnimation(self.fig, self._plot_path,


```




```

/tmp/ipykernel_3055857/955263672.py:37: UserWarning: frames=<generator object batched at 0x7...
    anim = animation.FuncAnimation(self.fig, self._plot_path,

```




```
/tmp/ipykernel_3055857/955263672.py:37: UserWarning: frames=<generator object batched at 0x7...
anim = animation.FuncAnimation(self.fig, self._plot_path,
```



```
maze = Maze8(maze_str)
success, search_path, node2parent, node2dist = astar(
    maze, partial(euclidean_heurist_dist, scale=1),
    start_pos, goal_pos)
#print(success, search_path)
assert success
anim = maze.animate(search_path, batch_size=20)
path = backtrace_path(node2parent, start_pos, goal_pos)
#maze.init_plots(reinit=True)
path_plot = maze.plot_path(path, color='k') # Draws the traced shortest path
anim
```

```
/tmp/ipykernel_3055857/955263672.py:37: UserWarning: frames=<generator object batched at 0x...
  anim = animation.FuncAnimation(self.fig, self._plot_path,
```




2.3.1 Admisibility and Consistency of heuristic function

1. An admissible heuristic is one that never overestimates the cost to reach the goal.
2. A heuristic $h(n)$ is consistent if it satisfies the triangle inequality:

$$h(s_t) \leq c(s_t, s_{t+1}) + h(s_{t+1}).$$

2.4 Converting maze to grid to graph



2.5 Other ways of converting a maze into a graph




Figure 3.31 A scene with polygonal obstacles. S and G are the start and goal states.

2.6 Rapidly exploring random trees

