

Discrete Planning

Vikas Dhiman


Friday 12th September, 2025

[Open in COLAB](#)

1 Required readings

- Ch 7 and 23, Carmen's Intro to Algorithms
- [Breadth-first search and its uses \(article\) | Khan Academy](#)
- [PythonRobotics/dijkstra.py](#)
- [Discrete Planning by Laval](#)
- 3.5.2 of Russel and Norvig: Artificial Intelligence

2 Discrete planning in Robotics



2.0.1 Abstraction of a planning problem

1. State space $s_t \in \mathcal{S}$. Property of a robot that changes with time and represents all the information needed for future planning (Markovian property).

sideways
t not Move t+1
t+2

For a robot that cannot move sideways and have to turn
state = (x, y, θ)
 $(-\pi, \pi)$

x, y, z 3D
 θ, ϕ, ψ roll/pitch/yaw
Euler angles

State = ? / State space
↳ a) Time varying property


b) Future action should be completely described by state

$a_t \perp s_{t-1} | s_t$
Markovian

$P(s_{t+1} | s_t, u_t) = ?$
state transition probabilities

system dynamics

$s_{t+1} = f(s_t, u_t)$
↑ next state ↑ state ↑ action



For example, 2D coordinate of a grid $s = (x, y)$. If orientation of a robot matters, for example, a car that can only move forward or backwards and needs to turn before it can move left or right, a state $s = (x, y, \theta)$. What would be state for a drone or a plane that can fly in 3D and rotate in 3D?

- ① State space: all possible states
- ② Action space: all possible action
- ③ $s_{t+1} = f(s_t, u_t)$
- ④ so
- ⑤ Goal state s_g

$s_t = \begin{bmatrix} x_t \\ y_t \end{bmatrix}$
 $u_t = \begin{bmatrix} \Delta x_t \\ \Delta y_t \end{bmatrix}$
down ↑ up
right ↑ left

2. Action space per state $u \in U(s)$. The choices of u can be made. For example, up, down, left right movement can be encoded as $U(s_t) = \{(0, -1), (0, 1), (1, 0), (-1, 0)\}$.

3. State transition function $s_{t+1} = f(s_t, u_t)$. For example, the up down left-right action can be combined as addition to get the next state $s_{t+1} = s_t + u_t$. When the system is stochastic, then it also known as state transition probabilities $P(s_{t+1} | s_t, u_t)$. In control systems, also known as system dynamics.

- 4. Initial State $s_I \in S$
- 5. Goal states $s_G \subseteq S$

$s_{t+1} = s_t + u_t$
 $f(s_t, u_t)$

2.0.2 A Graph

A graph $G = \{V, E\}$ is defined by a set of vertices V and a set of edges E such that each edge $e \in E$ is formed by a pair of start and end vertices $e = (v_s, v_e), v_s \in V, v_e \in V$. The first vertex is called the start of the edge $v_s = \text{start}(e)$ and second vertex is called the end $v_e = \text{end}(e)$.

A discrete planning problem can be converted into a graph by defining

- 1. Vertices as the state space V
- 2. The action space at each state as the edges connected to that vertex/state, $U(s_t) = \{(s_t, s_j) \mid (s_t, s_j) \in E\}$.
- 3. State transition function is the other end of th edge, $s_{t+1} = f(s_t, u_t) = \text{end}(u_t)$, where $s_t = \text{start}(u_t)$.

2.0.3 Representations of Graphs

(Chapter 23 of Introduction to Algorithms by Carmen et al)

Graph
Vertices/Nodes
Edge that connect them

$v_i \equiv s_i$
 $u_i \rightarrow (s_i, s_j)$

$s_{t+1} = \begin{cases} (x+1, y) & \text{if } u_t = R \\ (x-1, y) & \text{if } u_t = L \\ \vdots \end{cases}$

Discrete planning
↓
Graph
↓
Graph traversal problem

key →




Figure 23.1 Two representations of an undirected graph. (a) An undirected graph G having five vertices and seven edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .


Undirected graph

$\{$
 $1 : [(2, 2.0), (5, 3.0)]$
 $2 : [(1, 2.0), (3, 3.5), (4, 2.5), (5, 1.5)]$
 $\vdots$
 $\}$
 $O(|E|/|V|)$

$\begin{matrix} & 1 & 2 & 3 & 4 & 5 \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{pmatrix} 0 & 2.0 & 0 & 0 & 3.0 \\ & & & & \end{pmatrix} \end{matrix}$

$(2, 4)?$
 $O(1)$

Programmatically you can represent a adjacency list as python lists
 # Python lists are not linked lists, they are arrays under the hood
 G_adjacency_list = {
 1 : [2, 5], # node : list of neighbors
 2 : [1, 5, 3, 4],
 3 : [2, 4],
 4 : [2, 5, 3],
 5 : [4, 1, 2]
 }



$\xi =$ # Prefer to represent a matrix in python either as a list of lists or a numpy array
 import numpy as np
 G_adjacency_matrix = np.array([
 [0, 1, 0, 0, 1],
 [1, 0, 1, 1, 1],
 [0, 1, 0, 1, 0],
 [0, 1, 1, 0, 1],
 [1, 1, 0, 1, 0]
])

$O(1)$

$(1, 2) : 3.0$
 $(1, 5) : 4.5$
 $\{$

out neighbor start end nodes
 ↓
 nodes

Edge list is another possible representation
 G_edge_list = [
 (1, 2), (1, 5),
 (2, 1), (2, 5), (2, 3), (2, 4),
 (3, 2), (3, 4),
 (4, 2), (4, 5), (4, 3),
 (5, 4), (5, 1), (5, 2)
]

start end
 edge list = { (1, 2)
 (1, 4)
 ⋮
 }




Figure 23.2 Two representations of a directed graph. (a) A directed graph G having six vertices and eight edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .

Directed graph representation

Programmatically you can represent a adjacency list as python lists
 # Python lists are not linked lists, they are arrays under the hood.
 G_adjacency_list = {
 1 : [2, 4],
 2 : [5],
 3 : [6, 5],
 4 : [2],
 5 : [4],

```

    5 : [6]
}

# Prefer to represent a matrix in python either as a list of lists or a numpy array
import numpy as np
G_adjacency_matrix = np.array([
    [0, 1, 0, 1, 0, 0],
    [0, 0, 0, 0, 1, 0],
    [0, 0, 0, 0, 1, 1],
    [0, 1, 0, 0, 0, 0],
    [0, 0, 0, 1, 0, 0],
    [0, 0, 0, 0, 0, 1]
])

# Edge list is another possible representation
G_edge_list = [
    (1, 2), (1, 4),
    (2, 5),
    (3, 6), (3, 5),
    (4, 2),
    (5, 6)
]

# Exercise 1

# Write a function that converts a graph in adjacency list format to adjacency matrix and vice versa
def adjacency_list_to_matrix(G_adj_list):
    G_adj_mat = None # TODO: Write code to convert to adj_mat
    return G_adj_mat

def adjacency_matrix_to_list(G_adj_mat):
    G_adj_list = None # TODO: Write code to convert to adj_list
    return G_adj_list





# Use the above graphs to test
print(adjacency_list_to_matrix(G_adjacency_list))
print(adjacency_matrix_to_list(G_adjacency_matrix))

None
None


```

Tree
vs
2.0.4 Graph Search algorithms
Traversal

Trees don't have cycles/Loops



1. Breadth First Search ✓
2. Depth First Search ✓






Figure 23.3 The operation of BFS on an undirected graph. Tree edges are shown shaded as they are produced by BFS. Within each vertex u is shown $d[u]$. The queue Q is shown at the beginning of each iteration of the while loop of lines 9–18. Vertex distances are shown next to vertices in the queue.

Breadth first search (BFS)

```
from queue import Queue, LifoQueue, PriorityQueue
```

```
graph = {
    's' : ['w', 'r'],
    'r' : ['v'],
    'w' : ['t', 'x'],
    'x' : ['y'],
    't' : ['u'],
    'u' : ['y']
}

def bfs(graph, start, debug=False):
    seen = set() # Set for seen nodes (contains both frontier and dead states)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = Queue() # Frontier of unvisited nodes as FIFO
    node2dist = {start : 0} # Keep track of distances
    search_order = []
    seen.add(start)
    frontier.put(start)

    i = 0 # step number
    while not frontier.empty():
        # Creating loop to visit each node
        if debug: print("%d) Q = " % i, list(frontier.queue), end='; ')
        if debug: print("dists = " , [node2dist[n] for n in frontier.queue])
        m = frontier.get() # Get the oldest addition to frontier
        search_order.append(m)
```

```


    for neighbor in graph.get(m, []):
        if neighbor not in seen:
            seen.add(neighbor)
            frontier.put(neighbor)
            node2dist[neighbor] = node2dist[m] + 1
        else:
            assert node2dist[neighbor] <= node2dist[m] + 1, 'this should not happen in BFS'
            node2dist[neighbor] = min(node2dist[neighbor], node2dist[m] + 1)

    i += 1
    if debug: print("%d) Q = " % i, list(frontier.queue))
    return search_order, node2dist

print("Following is the Breadth -First Search order")
print(bfs(graph, 's', debug=True))    # function calling

Following is the Breadth -First Search order
0) Q = ['s']; dists = [0]
1) Q = ['w', 'r']; dists = [1, 1]
2) Q = ['r', 't', 'x']; dists = [1, 2, 2]
3) Q = ['t', 'x', 'v']; dists = [2, 2, 2]
4) Q = ['x', 'v', 'u']; dists = [2, 2, 3]
5) Q = ['v', 'u', 'y']; dists = [2, 3, 3]
6) Q = ['u', 'y']; dists = [3, 3]
7) Q = ['y']; dists = [3]
8) Q = []
(['s', 'w', 'r', 't', 'x', 'v', 'u', 'y'], {'s': 0, 'w': 1, 'r': 1, 't': 2, 'x': 2, 'v': 2,

```



BFS
Depth first search

gray = frontier
seen = black





Figure 23.4 The progress of the depth-first-search algorithm DFS on a directed graph. As edges are explored by the algorithm, they are shown as either shaded (if they are tree edges) or dashed (otherwise). Nontree edges are labeled B, C, or F according to whether they are back, cross, or forward edges. Vertices are timestamped by discovery time/finishing time.


```
graph = {
    's' : ['w', 'r'],
    'r' : ['v'],
    'w' : ['t', 'x'],
    'x' : ['y'],
    't' : ['u'],
    'u' : ['y']
}
```


DFS
Depth first
Search

BFS



seen = set({})
frontier = Queue()
DFS: FILO stack
BFS: FIFO
nodes that have been visited but their neighbors have not been




FIFO: First In First Out




find neighbors of w

seen
[s, w]
frontier (FILO)
[r, w]





find neighbors of x



seen
[s, w, t, x]
frontier
[r, t, x]
↑ intermediate

seen
[s, w, x, t]
frontier
[r, t, y]

BFS & DFS for shortest path



Breadth first search
give^a shortest path

for graphs with uniform
edge weights / distance

BFS → Dijkstra algorithm
(updates the visiting order based on
shortest distance)

FIFO
queue

Priority Queue


Out ←

4	3	3.5	2.1	2.0	1.0
v_1	v_2	v_3	v_4	v_5	v_6

← Priority

Highest priority node gets visited first


Dijkstra



seen

frontier

0
s




seen

s	u	x
---	---	---

frontier

-5	-10
x	u

Priority Queue



seen


s	x	u	v	y
---	---	---	---	---

frontier

-14	-7	-8
v	y	u

Priority Queue

-7	-8	-14
y	u	v



seen

s	x	u	v	y
---	---	---	---	---

frontier


Priority Queue

-8	-11	
u	v	

s	x	u	v	y
---	---	---	---	---

frontier

-9
x



computational complexity

LIFO Queue / stack (n)


FIFO Queue (n)

Priority Queue

enqueueing add $O(1)$ dequeueing remove $O(1)$



array



std::vector $O(1)$ $O(1)$


4	2	3	1	5
a	b	c	d	e

$O(1)$ $O(n)$

$O(n)$ $O(1)$

$O(\log(n))$ $O(1)$

heap data structure



Priority Queue → Priority Queue Updateable

}

```
def dfs(graph, start, debug=False):
    seen = set([start]) # List for seen nodes (contains both frontier and dead states)
    # Frontier is the boundary between seen and unseen (Also called the alive states)
    frontier = LifoQueue() # Frontier of unvisited nodes as FIFO
    node2dist = {start : 0} # Keep track of distances
    search_order = [] # Keep track of search order
    frontier.put(start)
```

frontier = Queue() →

```
    i = 0 # step number
    while not frontier.empty(): # Creating loop to visit each node
        if debug: print("%d Q = " % i, list(frontier.queue), end='; ')
        if debug: print("dists = " , [node2dist[n] for n in frontier.queue])
        m = frontier.get() # Get the oldest addition to frontier
        search_order.append(m)
```

```
        # adjacency list
        for neighbor in graph.get(m, []):
            if neighbor not in seen:
                seen.add(neighbor)
                frontier.put(neighbor)
                node2dist[neighbor] = node2dist[m] + 1
            else:
                node2dist[neighbor] = min(node2dist[neighbor], node2dist[m] + 1)
        i += 1
    if debug: print("%d Q = " % i, list(frontier.queue))
    return search_order, node2dist
```

graph[m] may throw Key Error
edge-distance

Driver Code

```
print("Following is the Depth -First Search path")
print(dfs(graph, 's', debug=True)) # function calling
```

Following is the Depth -First Search path

```
0) Q = ['s']; dists = [0]
1) Q = ['w', 'r']; dists = [1, 1]
2) Q = ['w', 'v']; dists = [1, 2]
3) Q = ['w']; dists = [1]
4) Q = ['t', 'x']; dists = [2, 2]
5) Q = ['t', 'y']; dists = [2, 3]
6) Q = ['t']; dists = [2]
7) Q = ['u']; dists = [3]
8) Q = []
(['s', 'r', 'v', 'w', 'x', 'y', 't', 'u'], {'s': 0, 'w': 1, 'r': 1, 'v': 2, 't': 2, 'x': 2,
```

Depth first search using recursion It is much easier to implement depth first search using recursion.

```
def dfs_recursive(graph, start,
                  seen=None, # List for seen nodes (contains both frontier and dead states)
                  search_order=None, # Keep track of search order
                  node2dist=None, # Keep track of distances
                  debug=False):
    if seen is None: seen = set()
    if search_order is None: search_order = []
    if node2dist is None: node2dist = {start: 0}


    seen.add(start) # List for seen nodes (contains both frontier and dead states)
    search_order.append(start)

    for neighbor in graph.get(start, []):
        edge_dist = 1 # for a weighted graph, you might want to get edge_weight from graph
        if neighbor not in seen:
            node2dist[neighbor] = node2dist[start] + edge_dist
            search_order, node2dist = dfs_recursive(
                graph, neighbor, seen=seen,
                search_order=search_order,
                debug=debug, node2dist=node2dist)
        else:
            node2dist[neighbor] = min(node2dist[neighbor], node2dist[start] + edge_dist)
    return search_order, node2dist



# Driver Code
print("Following is the Depth -First Search path")
print(dfs_recursive(graph, 's', debug=True)) # function calling
```

Following is the Depth -First Search path
 (['s', 'w', 't', 'u', 'y', 'x', 'r', 'v'], {'s': 0, 'w': 1, 't': 2, 'u': 3, 'y': 3, 'x': 2,


2.0.5 Search order in BFS vs DFS vs Dijkstra



Breadth first search vs Depth first search



Breadth first search vs Dijkstra



2.0.6 Converting a maze search to a graph search

```
import matplotlib.pyplot as plt
import numpy as np
maze_str = \
"""
+++++++
  +     +
+ + + +++
+ + +   +
+ + +   +

```

```

+ + +++ +
+      + +
+ +++ + +
+      +
+++++++
"""

```

```

class Maze:
    def __init__(self, maze_str, freepath=' '):
        self.maze_lines = [l for l in maze_str.split("\n")
                           if len(l)]
        self.FREEPATH = freepath
        self.fig = None

    def get(self, node, default):
        (r, c) = node
        m_row = self.maze_lines[r]
        nbrs = []
        if c - 1 >= 0 and m_row[c - 1] == self.FREEPATH:
            nbrs.append((r, c - 1))
        if c + 1 < len(m_row) and m_row[c + 1] == self.FREEPATH:
            nbrs.append((r, c + 1))
        if r - 1 >= 0 and self.maze_lines[r - 1][c] == self.FREEPATH:
            nbrs.append((r - 1, c))
        if r + 1 < len(self.maze_lines) and self.maze_lines[r + 1][c] == self.FREEPATH:
            nbrs.append((r + 1, c))
        return nbrs if len(nbrs) else default


    init_plots = init_plots
    plot_maze = plot_maze
    plot_step = plot_step
    plot_path = plot_path
    animate_search_path = animate_search_path

import matplotlib.pyplot as plt
import matplotlib.animation as animation
import matplotlib as mpl
%matplotlib inline
mpl.rc('animation', html='jshtml')

maze = Maze(maze_str)
search_path, node2dist = bfs(maze, (1, 0)) # prints the order of search all the searched nodes
maze.plot_maze()

maze.animate_search_path(search_path, node2dist)

```

```
def bfs_path(graph, start, goal):
    """
    Returns success and node2parent

    success: True if goal is found otherwise False
    node2parent: A dictionary that contains the nearest parent for node
    """
    seen = [start] # List for seen nodes.
    # Frontier is the boundary between seen and unseen
    frontier = Queue() # Frontier of unvisited nodes as FIFO
    node2parent = dict() # Keep track of nearest parent for each node (requires node to be h
    frontier.put(start)

    while not frontier.empty():
        # Creating loop to visit each node
        m = frontier.get() # Get the oldest addition to frontier
        if m == goal:
            return True, node2parent

        for neighbor in graph.get(m, []):
            if neighbor not in seen:
                seen.append(neighbor)
                frontier.put(neighbor)
                node2parent[neighbor] = m
```


```

    return False, []

def backtrace_path(node2parent, start, goal):
    c = goal
    r_path = [c]
    parent = node2parent.get(c, None)
    while parent != start:
        r_path.append(parent)
        c = parent
        parent = node2parent.get(c, None) # Keep getting the parent until you reach the start
        #print(parent)
    r_path.append(start)
    return reversed(r_path) # Reverses the path

maze = Maze(maze_str)
start = (1, 0)
goal = (8, 8)
success, node2parent = bfs_path(maze, (1, 0), (8, 8))
path = backtrace_path(node2parent, (1, 0), (8, 8))
#print(list(path))
maze.plot_path(path) # Draws all the searched nodes
plt.show()
#node2parent

```



2.1 Dijkstra algorithm




Figure 25.5 The execution of Dijkstra's algorithm. The source is the leftmost vertex. The shortest-path estimates are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. Black vertices are in the set S , and white vertices are in the priority queue $Q = V - S$. (a) The situation just before the first iteration of the **while** loop of lines 4–8. The shaded vertex has the minimum d value and is chosen as vertex u in line 5. (b)–(f) The situation after each successive iteration of the **while** loop. The shaded vertex in each part is chosen as vertex u in line 5 of the next iteration. The d and π values shown in part (f) are the final values.

2.1.1 PriorityQueue

PriorityQueue returns the smallest (or the largest) item in the queue faster than other data structures

```
import itertools
```

```
class MazeD(Maze):
```

```
    def get(self, node, default):
        nbrs = Maze.get(self, node, default)
        return zip(nbrs, itertools.repeat(1))
```

```
maze = MazeD(maze_str)
```

```
success, search_path, node2parent, node2dist = dijkstra(maze, (1, 0), (8, 8))
```

```
print(success, node2parent)
```

```
if success:
```

```
    path = backtrace_path(node2parent, (1, 0), (8, 8))
    maze.plot_path(path) # Draws all the searched nodes
```


```
True {(1, 0): None, (1, 1): (1, 0), (2, 1): (1, 1), (3, 1): (2, 1), (4, 1): (3, 1), (5, 1):
```

FIFO $\equiv$ BFS

LIFO $\Rightarrow$ DFS

Priority Queue $\equiv$ DIJKSTRA


- Shortest path
(uniform edge weights)
with +ve edge weights




```

maze = Maze8(maze_str)
success, search_path, node2parent, node2dist = dijkstra(maze, start_pos, goal_pos)
#print(success, search_path)
assert success
anim = maze.animate(search_path, title='Dijkstra')
path = backtrace_path(node2parent, start_pos, goal_pos)
#maze.init_plots(reinit=True)
path_plot = maze.plot_path(path, color='k') # Draws the traced shortest path
anim

```



```
path = backtrace_path(node2parent, (35, 9), (5, 50))
maze.init_plots(reinit=True, title='Dijkstra path')
maze.plot_path(path, color='r') # Draws the traced shortest path
plt.show()
```




2.2 Computational complexity notation

Asymptotically tight bound $\underline{\Theta}(g(n)) = \{f(n) : \text{there exist positive constants } c_1, c_2, \text{ and } n_o \text{ such that } 0 \leq c_1g(n) \leq f(n) \leq c_2g(n) \text{ for all } n \geq n_o\}$

Asymptotic upper bound $\underline{O}(g(n)) = \{f(n) : \text{there exist positive constants } c_2, \text{ and } n_o \text{ such that } 0 \leq f(n) \leq c_2g(n) \text{ for all } n \geq n_o\}$

Asymptotic lower bound $\Sigma(g(n)) = \{f(n) : \text{there exist positive constants } c_1, \text{ and } n_0 \text{ such that } 0 \leq c_1 g(n) \leq f(n) \text{ for all } n \geq n_0\}$



2.2.1 Computational complexity of BFS

Write down the computational complexity of each line in big-O notation $O()$

Assume the graph has $|V|$ nodes and $|E|$ edges

def bfs_barebones(graph, start):

seen = {start} # Set for seen nodes (contains both frontier and dead states) # $O(1)$
 # Frontier is the boundary between seen and unseen (Also called the alive states)
 frontier = Queue() # Frontier of unvisited nodes as FIFO # $O(1)$
 frontier.put(start) # $O(1)$

$O(|V|)$ while not frontier.empty(): # Creating loop to visit each node # $O(|V|)$
 m = frontier.get() # Get the oldest addition to frontier # $O(|V| * 1)$ $O(1)$
 for neighbor in graph.get(m, []): # $O(|V| * |E|/|V|) = O(|E|)$
 if neighbor not in seen: # $O(|E| * 1)$
 seen.add(neighbor) # $O(|E| * 1)$ - $O(1)$
 frontier.put(neighbor) # $O(|E| * 1)$

The computational complexity of BFS is $O(|E|)$. Some books write it as $O(|V| + |E|)$

where $O(|V|)$ is the cost of initializing states of different nodes

2.2.2 Computational complexity of Dijkstra

Write down the computational complexity of each line in big-O notation $O()$

Assume the graph has $|V|$ nodes and $|E|$ edges

def dijkstra_barebones(graph, start):

seen = {start} # Set for seen nodes (contains both frontier and dead states) # $O(1)$
 # Frontier is the boundary between seen and unseen (Also called the alive states)
 frontier = PriorityQueue() # Frontier of unvisited nodes as PriorityQueue # $O(1)$
 frontier.put((0, start)) # $O(1)$
 node2dist = {start: 0} # Keep track of cost to arrive at each node # $O(1)$

while not frontier.empty(): # Creating loop to visit each node
 dist_and_node = frontier.get() # Get the smallest dist node # $O(|V| \log(|V|))$
 m_dist = dist_and_node.dist
 m = dist_and_node.node

```

for neighbor, edge_dist in graph.get(m, []): #  $O(|V| * |E|/|V|) = O(|E|)$ 
    if neighbor not in seen: #  $O(|E| * 1)$ 
        seen.add(neighbor) #  $O(|E| * 1)$ 
        frontier.put(neighbor) #  $O(|E| * \log(1))$  # for fibonacci heap
        node2dist[neighbor] = m_dist + edge_dist #  $O(1)$ 
    elif node2dist[neighbor] > m_dist + edge_dist: #  $O(1)$ 
        node2dist[neighbor] = m_dist + edge_dist #  $O(1)$ 

```

The computational complexity of Dijkstra is $O(|V|\log(|V|) + |E|)$ when implemented
 # using a Fibonacci heap based PriorityQueue

2.3 PriorityQueue (Heaps Chapter 7 of Carmen's intro to algorithms)




Figure 7.1 A heap viewed as (a) a binary tree and (b) an array. The number within the circle at each node in the tree is the value stored at that node. The number next to a node is the corresponding index in the array.

Heap property

1. $H[\text{Parent}(i)] \geq H[i]$
 2. $\text{Parent}(i) = \text{ceil}(i/2)$
 3. $\text{LeftChild}(i) = 2i$
 4. $\text{RightChild}(i) = 2i+1$
- Handwritten notes:*
 - $\text{Parent}(i) = \text{ceil}(i/2)$ is labeled "1-based index"
 - $\text{LeftChild}(i) = 2i$ and $\text{RightChild}(i) = 2i+1$ are labeled "0 based index" and "1 based index" respectively.




Figure 7.2 The action of `HEAPIFY(A, 2)`, where `heap-size[A] = 10`. (a) The initial configuration of the heap, with `A[2]` at node $i = 2$ violating the heap property since it is not larger than both children. The heap property is restored for node 2 in (b) by exchanging `A[2]` with `A[4]`, which destroys the heap property for node 4. The recursive call `HEAPIFY(A, 4)` now sets $i = 4$. After swapping `A[4]` with `A[9]`, as shown in (c), node 4 is fixed up, and the recursive call `HEAPIFY(A, 9)` yields no further change to the data structure.

Heapify


HEAPIFY(A, i)

```

1  l ← LEFT(i)
2  r ← RIGHT(i)
3  if l ≤ heap-size[A] and A[l] > A[i]
4    then largest ← l
5  else largest ← i
6  if r ≤ heap-size[A] and A[r] > A[largest]
7    then largest ← r
8  if largest ≠ i
9    then exchange A[i] ↔ A[largest]
10   HEAPIFY(A, largest)

```

Heapify pseudocode




Heap Insert

Operation	$\overset{\text{max}}{\text{find-min}}$	$\overset{\text{max}}{\text{delete-min}}$	$\overset{\text{max}}{\text{insert}}$	$\overset{\text{max}}{\text{decrease-key}}$	$\overset{\text{max}}{\text{meld}}$
Binary ^[9]	$\Theta(1)$ ✓	$\Theta(\log n)$	$O(\log n)$	$O(\log n)$	$\Theta(n)$
Leftist	$\Theta(1)$	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)$	$\Theta(\log n)$
Binomial ^{[9][10]}	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)^{[a]}$	$\Theta(\log n)$	$O(\log n)$
Skew binomial ^[11]	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(\log n)$	$O(\log n)^{[b]}$
Pairing ^[12]	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$O(\log n)^{[a][c]}$	$\Theta(1)$
Rank-pairing ^[15]	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$\Theta(1)^{[a]}$	$\Theta(1)$
Fibonacci ^{[9][2]}	$\Theta(1)$	$O(\log n)^{[a]}$	$\Theta(1)$	$\Theta(1)^{[a]}$	$\Theta(1)$
Strict Fibonacci ^[16]	$\Theta(1)$	$O(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
Brodal ^{[17][d]}	$\Theta(1)$	$O(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
2-3 heap ^[19]	$O(\log n)$	$O(\log n)^{[a]}$	$O(\log n)^{[a]}$	$\Theta(1)$	?

Heap runtimes

2.4 A-star (A*) algorithm



(Required reading: 3.5.2 of Russel and Norvig: Artificial Intelligence)



```
from hw2_solution import PriorityQueueUpdatable
```

- ① BFS (even edge weights costs dist) $\xrightarrow{\text{FIFO Q}}$ Priority Q **DIJKSTRA**
- ② DFS (LIFO Q)

If you are exploring the entire space then DIJKSTRA is optimal



You can stop exploring when you have found The goal and return the shortest path

If we have some general idea about the direction of the goal, we direct our exploration towards the goal and find the goal sooner than Dijkstra.


A-star/A* is improvement over Dijkstra that uses a heuristic function $h(s)$, typically telling the shortest distance to the goal. $h(s) \leq d(s, g)$

Priority Q in Dijkstra, we used the accumulated distance so far $g(s_0, s_t)$



A* Priority Q, use $\underbrace{g(s_0, s_t) + h(s_t)}$ instead

$$\text{node2dist}[s_t] = g(s_0, s_t)$$

$$\text{node2parent}[s_t] = s_{t-1}$$



Heap-extract-max



Binary Heap

```

import sys
def astar(graph, heuristic_dist_fn, start, goal, debug=False, debugf=sys.stdout):
    """
    edgecost: cost of traversing each edge

    Returns success and node2parent


    success: True if goal is found otherwise False
    node2parent: A dictionary that contains the nearest parent for node
    """
    seen = set([start]) # Set for seen nodes.
    # Frontier is the boundary between seen and unseen
    frontier = PriorityQueueUpdatable() # Frontier of unvisited nodes as a Priority Queue
    node2parent = {start : None} # Keep track of nearest parent for each node (requires node2dist)
    hfn = heuristic_dist_fn # make the name shorter
    node2dist = {start: 0 } # Keep track of cost to arrive at each node  $g(s_0, s_t)$ 
    search_order = []
    frontier.put(PItem(0 + hfn(start, goal), start)) #  $g(s_0, s_t) + h(s_t, g)$  < - - - - - Different from dijkstra

    if debug: debugf.write("goal = " + str(goal) + '\n')
    i = 0
    while not frontier.empty():
        # Creating loop to visit each node
        dist_m = frontier.get() # Get the smallest addition to the frontier
        if debug: debugf.write("%d) Q = " % i + str(list(frontier.queue)) + '\n')
        if debug: debugf.write("%d) node = " % i + str(dist_m) + '\n')
        #if debug: print("dists = " , [node2dist[n.node] for n in frontier.queue])
        m = dist_m.node
        m_dist = node2dist[m]
        search_order.append(m)
        if goal is not None and m == goal:
            return True, search_order, node2parent, node2dist
        for neighbor, edge_cost in graph.get(m, []):
            old_dist = node2dist.get(neighbor, float("inf"))
            new_dist = edge_cost + m_dist # not using heuristic
            if neighbor not in seen:
                ✓seen.add(neighbor)  $g(s_0, s_t) + h(s_t, g)$ 
                ✓frontier.put(PItem(new_dist + hfn(neighbor, goal), neighbor)) # < - - - - -
                node2parent[neighbor] = m
                node2dist[neighbor] = new_dist
            elif new_dist < old_dist:
                node2parent[neighbor] = m
                node2dist[neighbor] = new_dist
            # ideally you would update the dist of this item in the priority queue
            # as well. But python priority queue does not support fast updates
            # - - - - - Different from dijkstra - - - - -

```

shortest distance
ignoring obstacles

$h(s_t) \approx d(s_t, g)$



A^* only
increases the
prob, that you
hit the goal first

Stop exploring
once you reach the
goal


```



        old_item = PItem(old_dist + hfn(neighbor, goal), neighbor)
        if old_item in frontier:
            frontier.replace(
                old_item,
                PItem(new_dist + hfn(neighbor, goal), neighbor))
        i += 1
    if goal is not None:
        return False, [], {}, node2dist
    else:
        return True, search_order, node2parent, node2dist

import math
from functools import partial


def euclidean_heurist_dist(node, goal, scale=1):
    x_n, y_n = node
    x_g, y_g = goal
    return scale*math.sqrt((x_n - x_g)**2 + (y_n - y_g)**2)

```






```
maze = Maze8(maze_str)
success, search_path, node2parent, node2dist = astar(
    maze, partial(euclidean_heurist_dist, scale=1),
    start_pos, goal_pos)
#print(success, search_path)
assert success
anim = maze.animate(search_path, batch_size=20, title='A -star heuristic=1.0')
path = backtrace_path(node2parent, start_pos, goal_pos)
#maze.init_plots(reinit=True)
path_plot = maze.plot_path(path, color='k') # Draws the traced shortest path
anim
```



2.4.1 Admissibility and Consistency of heuristic function

- Implies $\left\{ \begin{array}{l} 1. \text{ An admissible heuristic is one that never overestimates the cost to reach the goal.} \\ 2. \text{ A heuristic } h(n) \text{ is consistent if it satisfies the triangle inequality.} \end{array} \right. \quad h(s_t) \leq d(s_t, g)$
- $h(s_t) \leq c(s_t, s_{t+1}) + h(s_{t+1})$




Sum of any two sides $\geq$ Third side

$$h(s_t) \leq c(s_t, s_{t+1}) + h(s_{t+1})$$

A consistent heuristic is always admissible

An admissible but inconsistent²⁵ heuristic would require additional bookkeeping.

2.5 Converting maze to grid to graph



2.6 Path Planning in continuous spaces

Step 1: Converting a continuous space into a graph

Step 2: Plan the path on the graph using A-star

2.6.1 Converting a continuous space into a graph

For [convex polygon](#) obstacles, you can use the corners as nodes. Add an edge only if the straight line segment connecting the nodes is completely outside the obstacles.




Figure 3.31 A scene with polygonal obstacles. S and G are the start and goal states.